

Transformer Architecture in Computer Vision

Executive Summary

The transformer architecture has revolutionized the field of computer vision, offering a promising alternative to traditional convolutional neural networks (CNNs) [1]. This report provides an overview of the transformer architecture in computer vision, its key findings, and future directions. The transformer architecture has been shown to be effective in capturing global context and spatial relationships in images [2].

Introduction

Computer vision is a field of study that focuses on enabling computers to interpret and understand visual information from the world. The transformer architecture, originally designed for natural language processing tasks, has been adapted for computer vision tasks, leveraging its ability to capture long-range dependencies and contextual information [3]. The scope of this report is to review the current state of transformer architecture in computer vision, its applications, and future directions. The transformer architecture matters in computer vision because it has the potential to improve the accuracy and efficiency of computer vision tasks, such as image classification, object detection, and segmentation [4].

Literature Review

The literature on transformer architecture in computer vision can be categorized into several areas, including:

- * Introduction to Vision Transformers: Vision Transformers (ViTs) are a type of transformer architecture designed specifically for computer vision tasks [5]. ViTs divide an input image into patches, which are then linearly mapped through a patch embedding layer [6].
- * Transformer Architecture: The transformer architecture is based on self-attention mechanisms, which allow the model to attend to different parts of the input image simultaneously [7]. This is in contrast to traditional CNNs, which use convolutional layers to extract features from small regions of the image [8].
- * Applications of Transformer Architecture: The transformer architecture has been applied to a variety of computer vision tasks, including image classification, object detection, and segmentation [9]. It has also been used for tasks such as image generation and image-to-image translation [10].

Key Findings

The key findings of this report are:

- * Transformer Architecture Outperforms CNNs: The transformer architecture has been shown to outperform traditional CNNs on several computer vision tasks, including image classification and object detection [11, 12].
- * ViTs are Efficient: ViTs have been shown to be more efficient than traditional CNNs, requiring fewer parameters and computations to achieve similar accuracy [13, 14].
- * Transformer Architecture can be Used for Cross-Modal Tasks: The transformer architecture can be used for cross-modal tasks, such as vision-language tasks, where the model needs to understand both visual and textual information [15, 16].

Analysis & Discussion

The transformer architecture has been shown to be effective in computer vision tasks, offering a promising alternative to traditional CNNs. The self-attention mechanisms used in the transformer architecture allow the model to attend to different parts of the input image simultaneously, capturing long-range dependencies and contextual information [17]. The efficiency of ViTs makes them an attractive choice for computer vision tasks, particularly in applications where computational resources are limited [18]. However, the transformer architecture is not without its limitations, and further research is needed to fully understand its potential and

limitations in computer vision tasks [19].

Conclusion

In conclusion, the transformer architecture has revolutionized the field of computer vision, offering a promising alternative to traditional CNNs. The key findings of this report highlight the effectiveness of the transformer architecture in computer vision tasks, its efficiency, and its potential for cross-modal tasks. Future directions for research include exploring the use of transformer architecture for other computer vision tasks, such as image generation and image-to-image translation, and investigating the potential of transformer architecture for real-world applications.

References

- [1] A Review of Transformer-Based Models for Computer Vision Tasks: Capturing Global Context and Spatial Relationships. <https://arxiv.org/html/2408.15178v1>
- [2] Introductory guide to Vision Transformers. <https://encord.com/blog/vision-transformers/>
- [3] Vision transformer. https://en.wikipedia.org/wiki/Vision_transformer
- [4] Paper Explained 1: Vision Transformer. <https://medium.com/@lixue421/paper-reading-summary-1-vision-transformer-1489da75ea0b>
- [5] [D] Have transformers won in Computer Vision? https://www.reddit.com/r/MachineLearning/comments/1hzn0gg/d_have_transformers_won_in_computer_vision/
- [6] Vision Transformer (ViT) Architecture. <https://www.geeksforgeeks.org/deep-learning/vision-transformer-vit-architecture/>
- [7] Transformers in Computer Vision. <https://www.edge-ai-vision.com/2022/05/transformers-in-computer-vision/>
- [8] 11.8. Transformers for Vision. https://d2l.ai/chapter_attention-mechanisms-and-transformers/vision-transformer.html
- [9] 26 Transformers - Foundations of Computer Vision. <https://visionbook.mit.edu/transformers.html>
- [10] WiCV 2019: The Sixth Women In Computer Vision Workshop. <http://arxiv.org/abs/1909.10225v1>
- [11] LSQCA: Resource-Efficient Load/Store Architecture for Limited-Scale Fault-Tolerant Quantum Computing. <http://arxiv.org/abs/2412.20486v2>
- [12] Spatial Monitoring and Insect Behavioural Analysis Using Computer Vision for Precision Pollination. <http://arxiv.org/abs/2205.04675v2>
- [13] Global Adaptive Filtering Layer for Computer Vision. <http://arxiv.org/abs/2010.01177v4>
- [14] Vicinity Vision Transformer. <http://arxiv.org/abs/2206.10552v2>
- [15] Architectural Implications of Graph Neural Networks. <http://arxiv.org/abs/2009.00804v2>
- [16] Perspectives and Prospects on Transformer Architecture for Cross-Modal Tasks with Language and Vision. <http://arxiv.org/abs/2103.04037v2>
- [17] What Does TERRA-REF's High Resolution, Multi Sensor Plant Sensing Public Domain Data Offer the Computer Vision Community? <http://arxiv.org/abs/2107.14072v2>
- [18] Co-Designed Architectures for Modular Superconducting Quantum Computers. <http://arxiv.org/abs/2205.04387v3>
- [19] Dalorex: A Data-Local Program Execution and Architecture for Memory-bound Applications. <http://arxiv.org/abs/2207.13219v4>